

**From:** CHIIR 2026 Full and Perspective Papers chuir2026pc@easychair.org  
**Subject:** [CHIIR26] Notification for paper 2359  
**Date:** December 10, 2025 at 5:31 AM  
**To:** Divyesh Patel dpatel45@buffalo.edu



[You don't often get email from chuir2026pc@easychair.org. Learn why this is important at <https://aka.ms/LearnAboutSenderIdentification> ]

Dear Divyesh,

We regret to inform you that your submission titled:  
"Multimodal Knowledge Graph Representation with State-Space Models for Medical Image Retrieval" (2359)

has not been accepted for a FULL PAPER presentation at the 2026 ACM CHIIR conference.

We received 77 full/perspective paper submissions and accepted 27 (35%).

**PROCESS:** A minimum of three members of the Program Committee (PC) reviewed your paper. Then, a member of the Senior Program Committee (SPC) led an online discussion and wrote a summary meta-review. In the information below, you should be able to identify distinct reviews matching these different roles. In general, when reviewers had a difference of opinion after the online discussion, a final decision was reached by discussing the paper with the SPC member in charge.

**REVIEWS:** Reviews and meta-reviews for your paper are included at the bottom of this email. We hope that these reviews provide helpful feedback to help you improve your paper and resubmit it to ACM CHIIR or another similar conference or journal in the future.

**ATTENDANCE:** Thank you for your submission. We very much hope that you will still be able to join us for CHIIR 2026 in Seattle.

Further information about the conference (including registration information) is available at: <https://nam12.safelinks.protection.outlook.com/?url=https%3A%2F%2Fchuir2026.github.io%2F&data=05%7C02%7Cdpate145%40buffalo.edu%7C9ad338ff464c413a837708de37d70ff7%7C96464a8af8ed40b199e25f6b50a20250%7C0%7C639009595084690548%7CUnknown%7CTWFpbGZsb3d8eyJFbXB0eU1hcGkiOnRydWUsIIYIOilwLjAuMDAwMClsIIAiOjXaW4zMilsIkFOljoiTWFpbCIsIldUljoyfQ%3D%3D%7C60000%7C%7C%7C&sdata=PoQC02%2BbQTC9vJ9IM7GSw%2BbPPM53ZpiSOYah%2Fx6GJ0%3D&reserved=0>

---

#### STUDENT VOLUNTEERS

Student volunteers (SV) play a crucial role in ensuring that the conference runs smoothly and that attendees have an excellent experience. As an SV, your registration fee will be fully covered, and you will gain an opportunity to see how a major international research conference operates behind the scenes.

Apply to be an SV through <https://nam12.safelinks.protection.outlook.com/?url=https%3A%2F%2Fforms.gle%2FkbASXwKThYU6mQKs6&data=05%7C02%7Cdpate145%40buffalo.edu%7C9ad338ff464c413a837708de37d70ff7%7C96464a8af8ed40b199e25f6b50a20250%7C0%7C639009595084715351%7CUnknown%7CTWFpbGZsb3d8eyJFbXB0eU1hcGkiOnRydWUsIIYIOilwLjAuMDAwMClsIIAiOjXaW4zMilsIkFOljoiTWFpbCIsIldUljoyfQ%3D%3D%7C60000%7C%7C%7C&sdata=YLVDeGFOgAaQL3OXFc67wqwHaMZ4P0APiCBNeOPjRH%3D&reserved=0> by >>31 December 2025<< Notifications will go out by >>10 January 2026<<

---

We wish you all the very best with publishing your work in the future. Thank you for your hard work on this submission!

Best wishes,

Adam Fournay, Carla Teixeira Lopes, and Johanne Trippas  
ACM CHIIR 2026 Programme Chairs  
chuir26-program@acm.org

ACM CHIIR 2026 conference website: <https://nam12.safelinks.protection.outlook.com/?url=https%3A%2F%2Fchuir2026.github.io%2F&data=05%7C02%7Cdpate145%40buffalo.edu%7C9ad338ff464c413a837708de37d70ff7%7C96464a8af8ed40b199e25f6b50a20250%7C0%7C639009595084733024%7CUnknown%7CTWFpbGZsb3d8eyJFbXB0eU1hcGkiOnRydWUsIIYIOilwLjAuMDAwMClsIIAiOjXaW4zMilsIkFOljoiTWFpbCIsIldUljoyfQ%3D%3D%7C60000%7C%7C%7C&sdata=Vx6gDQ6rtcg5vKZLB3xBULFFAUhd4BNfachq6TVqmY%3D&reserved=0>

---

SUBMISSION: 2359

TITLE: Multimodal Knowledge Graph Representation with State-Space Models for Medical Image Retrieval

----- METAREVIEW -----  
---- This is the meta-review by the Senior PC member responsible for the paper ----

#### Synthesis

All three reviewers recommend weak reject and converge on similar assessments of the manuscript. Reviewers note that the paper proposes Graph CLIP State as a multimodal system that combines image text alignment, graph structure, and a temporal component to support interpretable medical image retrieval. Reported quantitative results are acknowledged as promising.

Major concerns raised by reviewers focus on evaluation and reporting. The dataset is small with four hundred cases and is not described in sufficient detail. Relevance grading and the use of graded metrics are found to be unclear and a custom metric that resembles DCG without normalization is not well justified. Reviewers also highlight presentation and clarity problems, noting that the manuscript is dense and difficult to follow in places and that experimental design and metric definitions require clearer explanation.

Scholarly grounding and artifact issues are also raised, including heavy reliance on arxiv preprint papers and reproduced figures without clear licensing or permission statements. Reviewer comments additionally emphasize the lack of user centered evaluation to support claims about

interpretability and trust, and gaps in discussion of scalability, privacy, deployment, and failure modes.

#### Final Recommendation

though the proposed system is conceptually interesting and the reported quantitative results are encouraging, the paper does not provide sufficient empirical rigor, clarity, or scholarly grounding for acceptance to the full paper track. Key issues include a small and under described dataset, unclear relevance criteria and metric justification, and presentation issues.

#### ----- REVIEW 1 -----

SUBMISSION: 2359

TITLE: Multimodal Knowledge Graph Representation with State-Space Models for Medical Image Retrieval

AUTHORS: Divyesh Patel, Piyush Gulhane, Pavithran Gnanasekaran, Rishab Darshan Shylendra and Alina Vereshchaka

##### ----- Overall evaluation -----

This study introduces Graph-CLIP-State (GCS), a novel multimodal retrieval-augmented reasoning framework designed for medical image retrieval. Research findings show that GCS achieves competitive retrieval performance (Precision@5 of 72.4% and MRR of 77.1%) and strong reliability (95.2% success rate) across neurological MRI datasets.

This study is limited by the design and evaluation of the experiment. A relatively small total dataset size (n=400), which was uniformly sampled and balanced. In addition, the evaluation relies on standard IR metrics (P@K, MRR, NDCG). While these are appropriate for retrieval quality, the authors correctly acknowledge that these metrics do not quantify interpretability or trust. A more user-centric evaluation of the proposed system would enhance the quality of this study.

Overall, this work bridges the divide between content-based medical image retrieval (CBMIR) and retrieval-augmented reasoning (RAR) by integrating multimodal embeddings (CLIP), graph-based knowledge (GraphRAG), and temporal State-Space Modeling (SSM). While the paper's goal is to enable interpretable, explainable, and user-centered retrieval systems, the research findings are primarily in technical achievement and quantitative IR metrics.

#### ----- REVIEW 2 -----

SUBMISSION: 2359

TITLE: Multimodal Knowledge Graph Representation with State-Space Models for Medical Image Retrieval

AUTHORS: Divyesh Patel, Piyush Gulhane, Pavithran Gnanasekaran, Rishab Darshan Shylendra and Alina Vereshchaka

##### ----- Overall evaluation -----

The manuscript is highly complex and dense with information. A clearer focus on the core contribution and a more streamlined presentation would greatly improve readability. At many points, the submission is difficult to follow. For example research design for the retrieval experiments also requires clearer description. The paper employs nDCG mostly used for multi-level evaluation (graded), yet it is not explained what constitutes relevance in this context, how many relevance levels are used, or on what criteria these levels were defined. In addition, the custom metric introduced in the paper corresponds essentially to DCG without normalization (i.e. not NDCG) its purpose and added value should be better justified. I recommend consulting and citing the foundational work by Järvelin & Kekäläinen (2002), which established DCG/nDCG and provides the conceptual framework necessary for interpreting graded relevance evaluations.

Moreover, the manuscript would benefit from a more balanced and rigorous engagement with the existing IR literature. A substantial portion of the references (17 items) are arXiv preprints. Because key components of the study build directly on these preprints, relying on non-reviewed versions weakens the scholarly foundation of the work. I encourage the authors to incorporate peer-reviewed sources whenever available and to more clearly situate the contribution within established information retrieval research.

Some figures reproduced from arXiv sources are included as direct copies without clarification of licensing or permissions. As preprints are not automatically released under open licenses, it would be helpful for the authors to ensure that permissions are properly addressed or to redraw the figures with appropriate attribution. Addressing these issues would significantly improve the clarity, reliability, and academic positioning of the manuscript.

#### ----- REVIEW 3 -----

SUBMISSION: 2359

TITLE: Multimodal Knowledge Graph Representation with State-Space Models for Medical Image Retrieval

AUTHORS: Divyesh Patel, Piyush Gulhane, Pavithran Gnanasekaran, Rishab Darshan Shylendra and Alina Vereshchaka

##### ----- Overall evaluation -----

---This review was an incorrect submission and has been removed by the PC Chairs // The scores of this review were not taken into account for the acceptance/rejection of the paper---

#### ----- REVIEW 4 -----

SUBMISSION: 2359

TITLE: Multimodal Knowledge Graph Representation with State-Space Models for Medical Image Retrieval

AUTHORS: Divyesh Patel, Piyush Gulhane, Pavithran Gnanasekaran, Rishab Darshan Shylendra and Alina Vereshchaka

##### ----- Overall evaluation -----

Summary

The paper presents Graph-CLIP-State (GCS), a system that groups medical images into a three-level graph so similar scans and their reports can be searched either as individual cases or as broader communities. It links images with text descriptions, builds stable clusters for reproducible results, and adds a time-aware component to summarize sequences of scans and provide traceable explanations. The goal is to make retrieval more interpretable and useful for clinical decision support.

Strength

-----

The approach feels practical and clinician-oriented by organizing cases into clear, explainable groups and showing the path that led to each result. Combining image–text alignment, deterministic clustering, and a temporal summarizer addresses real needs around trust, reproducibility, and longitudinal reasoning. The emphasis on traceable evidence and multiple retrieval modes (global, local, hybrid) strengthens its potential for real-world use.

#### Weakness

The paper lacks enough empirical detail to judge effectiveness: dataset sizes, quantitative results, and comparisons to strong baselines are not clearly reported. Questions remain about scalability, sensitivity to embedding quality, and how clinicians would evaluate or adopt the system in practice. Practical issues such as privacy, deployment constraints, and failure modes receive limited discussion.